Introducti on

NN (perceptron) depuis (60-70?) aujourd'hui, bc d'applications, CNN Notre but est alors de comprendre et savoir comment mettre en oeuvre, NN qui s'execute dans un temps raisonnable

Table of contents

1 Neural network concepts

A nerual netwok can be used for solving a wide array of problems in computer science, arguably the most common of which lies in computer vision. The purpose of this section is to explain and demonstrate the motivation behind the structure of a typical neural netoowork: the multi-layered perceptron (MLP). These neural networks can be applied to a wide range of problems, but for this section we are going to focus on computer vision.

The MLP performs well in image classification. The network can be considered a function with a number of inputs equal to the number of pixels in the image times three in the case of an RGB. The number of outputs is in this case the number of classes. This is obviously an extremely complicated function and finding the exact function is near impossible. A more intelligent approach is to approximate it with composition of several easily computed functions, such as linear functions. Figure 1 on page 3 presents a visualization of this concept.

The structured is loosely based on neuroscience; the nodes are often called neurons, which in reality is a number describing whether or not the neuron is active. The connections between neurons are called weights, denoted W, and translates how much the activation of one neuron should impact the activation of a following neuron. Lastly, we usually add constant values called biases noted b. These values can be seen as thresholds in order for a neuron to activate. Parallel neurons are joined in layers denoted by $X^{(l)}$ where l is the layer index.

An MLP is a so-called fully connected neural network. This means that each neuron in layer $X^{(l)}$, impacts the activation of each neuron in the following layer $X^{(l+1)}$. With these notations, the passage from layer to the next can be written
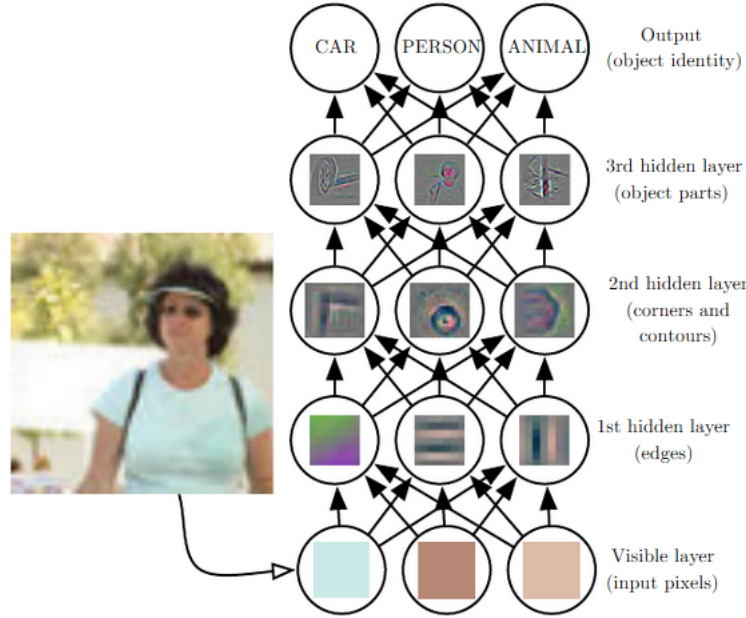
Figure 1 A visual

$$X^{(l+1)} = WX^{(l)} + b \tag{1}$$

A simple network might consist of two such layers with Y number of neurons each. This means (Y) number of adjustable parameters! With the structure defined, we have yet to determine the parameters. This is where the training phase enters.

1.1 Training phase

The idea is simple. We start out with a large dataset with n classified examples, called a training set and we adjust our network to fit these data. The hope is then that the network will be able to classify new images with these same parameters. In order to quantify the performance of the MLP during the training phase, we introduce a cost function that, in short, is a non-negative function that is closer to zero the better the network performs. Knowing the desired output with the corresponding input, we look to minimize this cost function by adjusting the parameters of the MLP i.e. its weights and biases. This can be seen as an optimization problem with thousands of inputs (in our case, Y weights and biases). For demonstration purposes, we will explain the concept of gradient descent, an intuitive and common optimization algorithm.

1.2 Optimization using gradient descent

We have a function that we want to minimize without calculating all its possible values. The gradient of the function is an operation that tells us which inputs to adjust and by how much, in order to increase the function by as much as possible in a given point. Visually, it can be considered the direction in the input space that increases the output the most. In our case, we want to decrease the cost by as much as possible so we take the negative value of the gradient.

As such, the gradient descent algorithm consists of, at each iteration, calculating the cost and its gradient, adjusting the inputs by "stepping in the direction of the gradient" (Figure 2) and repeating until we have found a local minimum or until the cost is sufficiently low. A limitation to this algorithm, is that there can be several local minima and there is no guarantee of finding the lowest one. In practice, though, the local minima in the context of neural networks are usually sufficiently low.

In the context of training a neural network, we have yet to introduce a way to calculate this very important gradient of the cost function. To do so, we first need to determine a mathematical expression of the neural network itself. In fact, the network typically consists of more than just linear functions.

1.3 Activation function

Up until now, we have introduced a transition from one layer to the next as a linear transformation. The problems we want the network to solve are usually far from linear, which is why we introduce a non-linear activation function phi. We now have

$$X^{(l+1)} = \phi(WX^{(l)} + b) \tag{2}$$

There are different variants of activation functions. In some cases, we want to avoid too large values. For optimization purposes, we also want the activation function to be derivable. Both the sigmoid function defined as

$$\phi(x) = \frac{1}{1 + e^{-x}} \tag{3}$$

and the arctan function have these desired properties, though the simple ReLU defined as

$$\phi(x) = \max(0, x)$$

is shown to be more efficient for optimization (source Goodfellow) and thus more common in practice.

1.4 Softmax

Before evaluating a cost function, we transform the output layer in a way that their values are positive and their sum is equal to 1. In this way, each value in the output layer can be interpreted

Figure 2 A visual

Figure 3 A visual

as the probability that a given input is classified as such. This is done using the softmax function defined as

$$\sigma(z_i) = \frac{\exp(z_i)}{\sum_{j=1}^{k} \exp(z_j)} \tag{4}$$

where k is equal to the number of outputs.

While in the training phase, we obviously know the desired output to each input, thus we want that specific neuron to have a value close to 1, and all the others to be close to 0. The loss function, which quantifies "how close" the network's guess is to a given output, can have different definitions depending on the problem. For instance, a common loss function that pairs well with the softmax function, is the negative log-likelihood defined as

$$L_i = -\log(y_i) \tag{5}$$

In order to quantify how well the network performs on the entire training set, we want to evaluate the loss function on each data-input. The cost function can then be defined as $\sum_{i=1}^{n} L_i$. Or the sum of losses.

1.5 Forward propagation

We now have an expression for each component in the neural network. Let's see how they are assembled in the forward propagation (Figure 4).

1.6 Backward propogation

Backward propagation, finding the gradient. As previously mentioned, the NN can be seen as a composition of functions of several functions of which we can easily find the derivate. Intuitively, in order to find the gradient of the cost function with respect to the weights and biases, we can just apply the chain rule (Figure 5). As we can see, the output layer is dependent on its preceding layer, weights and biases, which, in turn, depends on the previous parameters. This creates a cascade of gradients calculated from the output- to the input layer, hence the name backward propagation.

After iterating over the training set, we now have the gradients needed for the optimization algorithm, which in turn allows us to adjust the weights and biases in order to, iteration by iteration, reach a local minimum. In practice, this approach performs relatively well, though there are improvements to be made. In the following chapters, we will examine how convolutional neural networks differs from the classic MLP.

2 *Convolutional* Neural Networks

A convolutional neural network is an evolution of a classical multilayer percepetron network. Recall the basic principle underpinning how a normal neuron in a neural network is supposed to work. The neuron is supposed to 'look' for features in its input data. If the neuron 'thinks' that those features are present in the input data it 'fires'. Otherwise the neuron does not fire.

In a classical multilayer percepetron network this is implemented in the following way. Each neuron contains a vector of weights, a bias and an activation function. The input to the neuron—which must be a vector of equal length to the neuron's own weight vector—is combined with the neuron's weight vector using the dot product. The neuron's bias is added onto the result which in turn is passed to the activation function which finally determines if the neuron 'fires' or not.

Using several layers of neurons one can achive quite remarkable results using this implementation of a neural network. However, a multilayer percepetron is inherently limited. The major problem is that neurons in these kinds of networks only accept input that is in the form of a vector. This means that for applications where it is not natural for the input to be in a vector format, say image recognition, the input first has to be translated to a vector format. Usually this results in a loss of information contained in the input. In the typical case of image recongition, the input is in the form of one or more arrays of two dimentions. For a multilayer percepetron to treat this input, the images has to be 'flattend' into a vector of one dimention before it can be passed on to the network. This procedure eliminates some of the pixel relations in the image. To deduce this, consider the process of reconstructing a flattend image. If the image's dimentions prior to being flattend is not known, it is impossible, without the aid of pattern recognition, to reconstruct the image and be sure the reconstruction is equal to the original image.

To fix this problem, we can use a 'simple' solution. Instead of having the neuron contain a vector of weights, let it have an array of weights. If we change the neuron's vector of weights into an array, we also need to change the operation that is used to combine the weights with the input (which in a neuron of a multilayer percepetron is the dot product). There are two things to consider here. The purpose of the weights is to look for features or *patterns* in the input—by emphasizing or deemphasizing certain aspects—and the operation must reflect this purpose of the weights. Furthermore, the result of the operation should be a single number which, in a sense, representes the neurons 'initial' confidence that the feature it is looking for is present in the input. The operation which does both of these things is the *Hadamard product*. The Hadamard product can be viewed as an extension of the dot product to two dimentional arrays. It combines two arrays—of the same dimentions—by multiplying corresponding entries together and summing the results. Which is precisly what the dot product does with two vectors.

The Hadamard product (Hp) of two matricies A, B with entries $a_{i,j}, b_{i,j}$ of equal dimentions is

$$\text{Hp}(A, B) = \sum_j \sum_i a_{i,j} \cdot b_{i,j} \tag{6}$$

Another thing we need to take into account is the dimentions of our new weight array. If we were to proceed analogically to how a multilayer percepetron works, the array should have the same size as the input to the neuron. 'Connecting' each value in the input to an individual weight in the neuron. But this approach means that each neuron, in principle, looks for a single feature in the entire input at once. A more refined approach, is to let the neuron's weight array be smaller in size than its input. Instead of applying the weight array (using the Hadamard product) to the entire input at once, we apply it (still using the Hadamard product) to portions of the input seperately. Intuitivly, this means that the weight arrays are 'scanned' across the entire input image. This allows the neuron to be trained to look for a single feature, such as a 'sharp edge' or a 'round corner', in multiple areas of the input. The result of this method will no longer be a single number

representing how 'initialy' confident the neuron is that a particular feature is present in the input **as a whole**. Rather, the result becomes a *feature map*. Another two dimentional array which representes how 'initialy' confident the neuron is that a particular feature is present in **specific locations** of the input.

Going one step further, we can allow the neuron to treat inputs of not only a single two dimentional array, but several two dimentional arrays. A typical example where the input would consists of several interlinked two dimentional arrays is RGB images. An RGB image consists of three arrays of pixel values (numbers) that desribe how red, green and blue an image is in each pixel. The number of interlinked two dimentional arrays present in the input, is known as the number of *channels* that the input has. In order for our neuron to treat inputs with more than one channel we let the neuron have as many channels as the input. That is to say, we equip the neuron with a weight array for each channel in the input. For each channel the weight arrays are applied to the input (using Hadamard) and the result in each channel is combined to form a final single feature map.

A neural network which makes use of layers of neurons of this kind, is a convolutional neural network. The multi-channel two dimentional arrays of weights inside each such neuron is called the neuron's *kernel* or *filter*. Why are these neural networks called convolutional neural networks? That will be explained in the next section.

2.1 The convolutional neuron

Let us start this section with a simple convolutional neuron. The neuron's kernel consists of a single weight matrix (one channel), some bias and some activation function. Consequently, the input that this neuron accepts is any one channel two dimentional array of size greater than its kernel.

The process of calcuating the neurons 'initial' confidence that the feature it is looking for is present in the input, is illustrated by Figure 4 on page 9. Figure 4 shows how the kernel 'scans' the entire input array to compute the feature map. Let us construct the general formula for this feature map.

We will denote the feature map, kernel and input as F, K and M respectivly. They are two dimentional arrays and their individual entries will be denoted as $F(x, y)$ where x is the column index and y the row index. The indexes will all start at zero (e.g. $M(1, 2)$ is the entry in the second column, third row of M). There are many lengths involved in working with three seperate arrays, various lengths will here be denoted by n. A subscript, either x or y, will indicate if it is a horisontal or vertical length and a subsequent argument will indicate which array the length belongs to. So for example, $n_x(M)$ is the horisontal length of the input array—the number of columns or x's in M.

If the input represents a typical MNIST image with dimentions 28x28, then $n_x(M) = 28$ and $n_y(M) = 28$. In this scenario, the indices of M range in $[0 \,.. \, n_x(M) - 1] \times [0 \,.. \, n_y(M) - 1]$.

Let us say that the kernel scans its input one column and one row at a time. In this scenario, the output feature map will have dimentions

2.2 The convolutional layer

Okay this is where I will start developing the mathematics. First we need to define the various variables used and some terminology. A network has several layers and each layer has its own specific attributes. A supersricpt enclosed in parantheses will be used as the layer index.

$$\delta x_Z^{(l+1)} = \delta x_W^{(l)} - \delta x_Z^{(l)} + 1$$

$$y_Z^{(l+1)} = \delta y_W^{(l)} - \delta y_Z^{(l)} + 1 \tag{7}$$

$$x \in [0 \,..\, n_x(F) - 1] \qquad y \in [0 \,..\, n_y(F) - 1]$$

$$x \in [0 \,..\, n_x(K) - 1] \qquad y \in [0 \,..\, n_y(K) - 1] \tag{8}$$

The product in layer l before activation, here denoted Z is given by the following formula
Regular formula

$$Z_{x,y,c}^{(l+1)} = \sum_{c'=0}^{\delta c_Z^{(l)}} \sum_{y'=0}^{\delta y_Z^{(l)}} \sum_{x'=0}^{\delta x_Z^{(l)}} \left(W_{x',y',c'}^{(l,c)} Z_{x+x',y+y',c'}^{(l)} \right) + b^{(l,c)} \tag{9}$$

Derivative with respect to Z. $c^* \in [c \,..\, c + \delta c_W^{(l)}]$. $y^* \in [y \,..\, y + \delta y_W^{(l)}]$. $x^* \in [x \,..\, x + \delta x_W^{(l)}]$.

$$
\begin{aligned}
\frac{\partial Z_{x,y,c}^{(l+1)}}{\partial Z_{x^*,y^*,c^*}^{(l)}} &= \frac{\partial}{\partial Z_{x^*,y^*,c^*}^{(l)}} \left(\sum_{c'=0}^{\delta c_W^{(l)}} \sum_{y'=0}^{\delta y_W^{(l)}} \sum_{x'=0}^{\delta x_W^{(l)}} \left(W_{x',y',c'}^{(l,c)} Z_{x+x',y+y',c'}^{(l)} \right) + b^{(l,c)} \right) \\
&= \sum_{c'=0}^{\delta c_W^{(l)}} \sum_{y'=0}^{\delta y_W^{(l)}} \sum_{x'=0}^{\delta x_W^{(l)}} \frac{\partial}{\partial Z_{x^*,y^*,c^*}^{(l)}} \left(\left(W_{x',y',c'}^{(l,c)} Z_{x+x',y+y',c'}^{(l)} \right) + b^{(l,c)} \right) \\
&= W_{x',y',c'}^{(l,c)}
\end{aligned}
\tag{10}
$$

Derivative with respecto to W. $c^* \in [0 \,..\, \delta c_W^{(l)}]$. $y^* \in [0 \,..\, \delta y_W^{(l)}]$. $x^* \in [0 \,..\, \delta x_W^{(l)}]$.

$$
\begin{aligned}
\frac{\partial Z_{x,y,c}^{(l+1)}}{\partial W_{x^*,y^*,c^*}^{(l,c)}} &= \frac{\partial}{\partial W_{x^*,y^*,c^*}^{(l,c)}} \left(\sum_{c'=0}^{\delta c_W^{(l)}} \sum_{y'=0}^{\delta y_W^{(l)}} \sum_{x'=0}^{\delta x_W^{(l)}} \left(W_{x',y',c'}^{(l,c)} Z_{x+x',y+y',c'}^{(l)} \right) + b^{(l,c)} \right) \\
&= \sum_{c'=0}^{\delta c_W^{(l)}} \sum_{y'=0}^{\delta y_W^{(l)}} \sum_{x'=0}^{\delta x_W^{(l)}} \frac{\partial}{\partial W_{x^*,y^*,c^*}^{(l,c)}} \left(\left(W_{x',y',c'}^{(l,c)} Z_{x+x',y+y',c'}^{(l)} \right) + b^{(l,c)} \right) \\
&= Z_{x^*+x',y^*+y',c^*}^{(l)}
\end{aligned}
\tag{11}
$$

Derivative with respect to b

$$
\begin{aligned}
\frac{\partial Z_{x,y,c}^{(l+1)}}{\partial b^{(l,c)}} &= \frac{\partial}{\partial b^{(l,c)}} \left(\sum_{c'=0}^{\delta c_W^{(l)}} \sum_{y'=0}^{\delta y_W^{(l)}} \sum_{x'=0}^{\delta x_W^{(l)}} \left(W_{x',y',c'}^{(l,c)} Z_{x+x',y+y',c'}^{(l)} \right) + b^{(l,c)} \right) \\
&= 1
\end{aligned}
\tag{12}
$$

If we add a bias b to this filter, the formula becomes

$$F_*(x,y) = \sum_y \sum_x \left(K(x,y) \cdot M(x+x, y+y) + b \right) \tag{13}$$

Sending this formula throug an activation function such as ReLU or the sigmoid function, we obtain the final output feature.

$$F_*(x,y) = \text{activation} \left(\sum_y \sum_x \left(K(x,y) \cdot M(x+x, y+y) + b \right) \right) \tag{14}$$

Figure 4 The basic forward operation of a convolutional layer

2.2.1 Where is the convolution?

In mathematics, the convolution operation is denoted by $*$ and the convolution s (itself a function) of two functions k and m is defined as

$$s(t) = (k * m)(t) = \int\limits_{-\infty}^{+\infty} k(x) \cdot m(t - x)\, \mathrm{d}x \tag{15}$$

$$s(x) = (k * m)(x) = \sum_{x=-\infty}^{+\infty} k(x) \cdot m(x - x) \tag{16}$$

For continuous (Eq. (15)) and discrete (Eq. (16)) functions respectivly. On a computer, we are in practice always working with discrete convolutions so for us the convolution of interest is Eq. (16). If we make the functions k and m two dimentional—they take two arguments as their input—their convolution also becomes two dimentional.

$$s(x, y) = (k * m)(x, y) = \sum_{y=-\infty}^{+\infty} \sum_{x=-\infty}^{+\infty} k(x, y) \cdot m(x - x, y - y) \tag{17}$$

Let us known say that k and m are two functions which each index an array of two dimentions. Meaning that the functions take two indexes as their input arguments. Let us further assume that k and m's arrays are zero at any index but for those contained in a small area. For k that area is

9

$[0 .. n_x(k) - 1] \times [0 .. n_y(k) - 1]$ and for m it is $[0 .. n_x(m) - 1] \times [0 .. n_y(m) - 1]$. In this case, the convolution of k and m reduces to

$$s(x,y) = (k * m)(x,y) = \sum_{y=0}^{n_y(k)-1} \sum_{x=0}^{n_x(k)-1} k(x,y) \cdot m(x-x, y-y) \tag{18}$$

Where $x \in [0 .. n_x(m) - 1]$ and $y \in [0 .. n_y(m) - 1]$.

The resemblance to Eq. (9) is already apparent but there is a subtle difference. Eq. (18) subtracts the blue indices from the red ones while Eq. (9) adds them together. The effect of subtracting instead of adding renders the convolution operation $(*)$ commutative.

$$(k * m)(x,y) = (m * k)(x,y)$$
$$\sum_x \sum_y k(x,y) \cdot m(x-x, y-y) = \sum_x \sum_y m(x,y) \cdot k(x-x, y-y) \tag{19}$$

So what is the effect of replacing the pluses in Eq. (9) with minuses? Well at first glance it seems that you cannot do it as you would end up with indexes that are out of bounds for the input image M. But let us assume that M behaves as m meaning that it is zero everywhere where the indices are out of bounds. What you find if you visualize the computation of F using Eq. (18), is that it is effectivly the same as using Eq. (9) with the kernel flipped about its horisontal and vertical axes. Which in turn means that flipping the kernel relative to the input, renders the computation of the feature map commutative. But this is a property that is usefull in mathematical contexts and less so for computing neural networks. It is Eq. (9) that is commonly used in macine learning applications and the associated mathematical operations is actually what is called the *cross-correlation*. Despite this, neural networks that implement these kinds of neurons are called convolutional even if the underlying mathematical operations is more often than not a close relative of the proper convolution operation.

2.2.2 Multiple channels

Let us return to Eq. (14) for a minute and ask: what happens if we add some channels to the input? Let us denote the number of channels by n_c. The input and kernel are now multi-channel arrays, their individual entries will be denoted as $K(i, j, c)$ with c ranging in $[0 .. n_c - 1]$ To calculate the feature map when multiple channels are present, we have to take into account the contribution of each individual channel. We do this by simply adding them together. Since we only want to add the bias of the filter once per entry in the feature map, the resulting modified formula becomes

$$F(x,y) = \text{activation} \left(\sum_y \sum_x \left(\sum_c \Big(K(x,y,c) \cdot M(x+x, y+y, c) \Big) + b \right) \right) \tag{20}$$

2.2.3 The stride

So far we have considered a feature map produced by scanning the kernel over the input one *step* at a time. As is demonstrated in Figure 4 on page 9. However, we could let the kernel take longer steps instead of moving just one row or one column at a time. The 'length' of each step that the kernel takes while moving over the input is what is known as the *stride*. Which we can further divide into the horisontal stride s_x and vertical stride s_y.

What purpose does altering the stride serve? Increasing the stride reduces the size of the resulting feature map since it is computed at fewer areas. Which has the net effect that the

computational complexity of the neuron is reduced. However, it also means that the neuron studies its input less which can result in vital information in the input being overlooked. This is not a big problem as long as the features of the input is judged to be larger than the kernel that the neuron uses. So long as this is the case, increasing the stride of a neuron can significantly reduce its computational complexity without resulting in it no longer surving its purpose.

Let us now further develop Eq. (20) to take into account variable strides. It is as simple as multiplying the red indices with the strides.

$$F(x,y) = \text{activation} \left(\sum_y \sum_x \left(\sum_c \Big(K(x,y,c) \cdot M(xs_x + x, ys_y + y, c) \Big) + b \right) \right) \tag{21}$$

2.2.4 Zero padding

If we increase a neurons stride we can shrink the size of the resulting feature map. But it does not allow us to increase the size of the feature map! To do this, a techinque known as *zero padding* is used. Instead of working with the input *as is*, one or more extra outer rows and columns of zeros are added onto it before the kernel is scanned across it. This allows us to increase the size of the resulting feature map.

This techinque is predominantly used to make the feature map have the same dimentions as the input image in applicatons where this is desirable.

2.3 The pooling 'neuron'

Downsampling and purpose of downsampling. Noise reduction and computational reduction.

2.4 Backward Propagation

Obtain formula for derivative of convolution and pooling.

3 Python Implementation

Knwon that we have taken a look at the theory behind convolutional neural networks, let us look at an implementation of one. The nerual network we are going to be looking at is one written by Alejandro Escontrela. It is written in Python using the NumPy library and is publicly avaliable at `https://github.com/Alescontrela/Numpy-CNN.git`

The network is going to tackle the classic neural network problem. Categorizing the handwritten digits in the MNIST database. Compared to other problems convolutional neural networks are commonly faced with, categorizing the MNIST images is rather simple. So the network will use a comperativly simple architecture as shown in Figure 5 on page 12. The network uses no zero padding. The first two layer consists of two indentical convolutional layers. They are in turn follwed by a single max-polling layer which notably, uses a stride of 2. The output is then flattened and passed to a dense layer with 128 neurons. The output of this layer is passed to the final dense layer, which necessarily consists of 10 neurons.

3.1 Convolution function

Here is the python code for the forward convolution operation of the network

Figure 5 The architecture of the network

```python
def convolution(image, filt, bias, s=1):
    '''
    Confolves `filt` over `image` using stride `s`
    '''
    (n_f, n_c_f, f, _) = filt.shape # filter dimensions
    n_c, in_dim, _ = image.shape # image dimensions

    out_dim = int((in_dim - f)/s)+1 # calculate output dimensions

    assert n_c == n_c_f, "Dimensions of filter must match dimensions of input image"

    out = np.zeros((n_f,out_dim,out_dim))

    # convolve the filter over every part of the image, adding the bias at each step.
    for curr_f in range(n_f):
        curr_y = out_y = 0
        while curr_y + f <= in_dim:
            curr_x = out_x = 0
            while curr_x + f <= in_dim:
                out[curr_f, out_y, out_x] = np.sum(filt[curr_f] * image[:,curr_y:curr_y+f,
curr_x:curr_x+f]) + bias[curr_f]
                curr_x += s
                out_x += 1
            curr_y += s
            out_y += 1

    return out
```

3.2 Max pooling

Here is the python code for the forward max-pooling operation of the network

3.3 Various functions

The network uses ReLU for its activation function,

```python
def maxpool(image, f=2, s=2):
    '''
    Downsample `image` using kernel size `f` and stride `s`
    '''
    n_c, h_prev, w_prev = image.shape

    h = int((h_prev - f)/s)+1
    w = int((w_prev - f)/s)+1

    downsampled = np.zeros((n_c, h, w))
    for i in range(n_c):
        # slide maxpool window over each part of the image and assign the max value
at each step to the output
        curr_y = out_y = 0
        while curr_y + f <= h_prev:
            curr_x = out_x = 0
            while curr_x + f <= w_prev:
                downsampled[i, out_y, out_x] = np.max(image[i, curr_y:curr_y+f, curr_x:curr_x+f])
                curr_x += s
                out_x += 1
            curr_y += s
            out_y += 1
    return downsampled

def convolutionBackward(dconv_prev, conv_in, filt, s):
    '''
    Backpropagation through a convolutional layer.
    '''
    (n_f, n_c, f, _) = filt.shape
    (_, orig_dim, _) = conv_in.shape
    ## initialize derivatives
    dout = np.zeros(conv_in.shape)
    dfilt = np.zeros(filt.shape)
    dbias = np.zeros((n_f,1))
    for curr_f in range(n_f):
        # loop through all filters
        curr_y = out_y = 0
        while curr_y + f <= orig_dim:
            curr_x = out_x = 0
            while curr_x + f <= orig_dim:
                # loss gradient of filter (used to update the filter)
                dfilt[curr_f] += dconv_prev[curr_f, out_y, out_x] * conv_in[:, curr_y:curr_y+f,
curr_x:curr_x+f]
                # loss gradient of the input to the convolution operation (conv1 in
the case of this network)
```

```
                dout[:, curr_y:curr_y+f, curr_x:curr_x+f] += dconv_prev[curr_f, out_y,
out_x] * filt[curr_f]
                curr_x += s
                out_x += 1
            curr_y += s
            out_y += 1
        # loss gradient of the bias
        dbias[curr_f] = np.sum(dconv_prev[curr_f])

    return dout, dfilt, dbias

def maxpoolBackward(dpool, orig, f, s):
    '''
    Backpropagation through a maxpooling layer. The gradients are passed through the
indices of greatest value in the original maxpooling during the forward step.
    '''
    (n_c, orig_dim, _) = orig.shape

    dout = np.zeros(orig.shape)

    for curr_c in range(n_c):
        curr_y = out_y = 0
        while curr_y + f <= orig_dim:
            curr_x = out_x = 0
            while curr_x + f <= orig_dim:
                # obtain index of largest value in input for current window
                (a, b) = nanargmax(orig[curr_c, curr_y:curr_y+f, curr_x:curr_x+f])
                dout[curr_c, curr_y+a, curr_x+b] = dpool[curr_c, out_y, out_x]

                curr_x += s
                out_x += 1
            curr_y += s
            out_y += 1

    return dout
```

- The problem the network is going to tackle (classic MNIST image recognition)

- The architecture of the network

- Choice of various functions and their associated code

- Results with time performance

- Possible Modifications

4 Conclusion

The advantages of a CNN over a NN.